

Personal Summary

Data Engineer with 2+ years of experience in healthcare data engineering, business intelligence, and analytics. Skilled in building scalable data pipelines and reporting solutions using Databricks, SQL, PySpark, MicroStrategy, and Power BI. Experienced across the full data lifecycle, including ingestion, transformation, modeling, validation, and delivery. Also proficient in GitHub, Azure DevOps, CI/CD, Liquibase, Databricks Asset Bundles, and Databricks CLI, with a strong understanding of data governance, security, and HIPAA compliant reporting for US based healthcare clients. Comfortable working with large and complex datasets, troubleshooting data issues, and improving data quality and reliability. Focused on building efficient data solutions that support both business reporting and day to day healthcare operations.

Skills :

- **Databases:** MySQL, PostgreSQL
- **BI & Analytics:** MicroStrategy, Power BI
- **Cloud & Data Platforms:** Azure Databricks, Databricks Asset Bundles (DABs), Databricks CLI
- **DevOps & Deployment:** Git, GitHub, Azure DevOps, CI/CD, Docker, Gitlab
- **Data Governance & Security:** Data Validation, Data Governance, Row-Level Security (RLS), PHI Controls, HIPAA
- **Programming:** Python, PySpark, SQL
- **AI & Productivity Tools:** GitHub Copilot, Claude, OpenAI API

Work Experience :

Data Engineer, Neurealm (Previously GAVS Technologies) | Jul 2024 – Present

I work in the healthcare domain, building data pipelines and reporting solutions for US based hospital clients. My role has grown beyond dashboard development, and I now work across the full data lifecycle, from ingestion and transformation to validation and delivery, with a focus on keeping the data accurate, reliable, and useful.

- Build and maintain data pipelines and data transformation workflows for US based healthcare clients, supporting operational reporting, analytics, and clinical decision making.
- Develop complex SQL and PySpark workflows in Azure Databricks to cleanse, transform, validate, and analyze large healthcare datasets for reporting and business requirements.
- Maintain documentation of critical metrics, data sources, and data mapping, so that definitions stay consistent and anyone picking up the work later understands exactly how a number is calculated.
- Supervise data quality through ongoing validation, catching inconsistencies or anomalies in large healthcare datasets before they reach a client facing report.
- Develop and maintain MicroStrategy dashboards, reports, Intelligent Cubes, advanced metrics, prompts, filters, attributes, and custom groups to deliver operational and clinical KPIs.
- Coordinate with other team members on schema changes and deployments using Liquibase, tracking releases across dev, QA, demo, and production so nothing breaks silently.
- Led the migration of 2TB+ of data from Teradata to Azure Databricks, ensuring data accuracy, validation, and reporting continuity throughout the migration.

- Manage schema changes and database deployments using Liquibase, coordinating releases across Dev, QA, Demo, and Production environments.
- Use SQL performance tuning and optimization regularly to keep reporting systems fast and reliable, especially as data volume grows.
- Apply Row-Level Security (RLS), object security, data anonymization, and PHI access controls to support secure and HIPAA compliant healthcare reporting.
- Performed SQL query troubleshooting and performance optimization, improving reporting efficiency and reducing query execution time in Databricks.
- Implement Databricks Asset Bundles (DABs) and Databricks CLI within CI/CD workflows to deploy schema and view changes and manage database scripts across multiple environments.
- Perform data profiling, quality validation, regression testing, and anomaly investigation to ensure accuracy and consistency across reporting systems for UAT and internal QA testing.
- Design and maintain data models and curated datasets in Azure Databricks, structuring transformed healthcare data into reliable and reusable data layers for downstream reporting, analytics, and business intelligence.
- Collaborate with client BA's, technical, and client stakeholders to gather data and reporting requirements, translate business needs into technical solutions, and resolve the incidents.
- Work within Agile Scrum teams, using Azure Boards for sprint management, GitHub for version control, and CI/CD pipelines for deployments across multiple environments.

Certifications :

Duration	Course	Institution
Aug 26 - Sep 26	Databricks Certified Data Engineer Professional	Databricks
Sep 25 - Dec 25	Data Science and Machine Learning	Gilbert Research Center
July 2025	Data Analysis and Visualization in Power BI	Coursera
Dec 2024	Generative AI for everyone	Coursera
Nov 2024	Microstrategy for Business Intelligence	Udemy
Aug 23 – Dec 23	Machine Learning using Python	Infosys Springboard

Education :

Madras Institute of Technology, Anna University Chennai, India

Bachelor of Technology in Electronics and Communication

Nov 2020 - M a y 2024

- Secured an aggregate of 7.7 CGPA

Higher Secondary Education, Vels Vidyashram Chennai, India

June 2018 – April 2020

Percentage-86.7%

Secondary Education, Vels Vidyashram Chennai, India

July 2005 – April 2018

Percentage-72.8%

Additional information :

- Co-authored and published a research paper on Machine Learning-based authentication using ECG and PPG signals and developed IoT-based surveillance systems for flood monitoring using computer vision.
- Completed an internship as an AI Developer at Truetech Solutions (July–August 2022), where I developed a face recognition-based interview chatbot using Flask, OpenCV, and Dlib to enable real-time candidate identity verification and AI-driven authentication.
- Languages Known: English, Tamil, and German.